

Bayesian imaging with deep generative priors

Marcelo Pereyra

Heriot-Watt University & Maxwell Institute for Mathematical Sciences

Joint work with Alessio Spagnolletti, Jean Prost, Andres Almansa and Nicolas Papadakis




24 June 2025, Glasgow
30th Biennial Numerical Analysis Conference

Background

Problem:

- Image data often not useful in raw form (limited resolution & noisy, or accurate but too expensive).
- Evidence-based decision-making needs accurate solutions and reliable uncertainty quantification.



Vision:

Use mathematics to upgrade imaging instruments into smart decision-making support systems.

Approach:

A probabilistic computational imaging framework integrating physical and generative AI models, Bayesian statistical decision-theory and fast (exa)scalable stochastic algorithms.

Today's talk: Generative AI-based Bayesian Imaging


Example of image generated by a Vision Language Model (VLM). These are probabilistic generative models represented by massive deep neural nets.



Prompt: Beautiful white Mediterranean outdoor courtyard, decorated with string lights and candles... Credit: Midjourney.com

Extreme noise Extreme Zoom

Key breakthrough: new mathematical underpinning allows embedding physical models into VLMs and *prompting* with physical measurements, while self-adjusting text prompts.



Problem Statement

We are interested in an unknown image $x^* \in \mathbb{R}^d$

We measure $y = Ax^* + w$

Recovering x^* from y is not well posed.


Bayesian Statistical Framework

Model x^* as a realisation of \tilde{x} and y as a realisation of $(y|\tilde{x} = x^*)$

We draw inferences about \tilde{x} having observed $\tilde{y} = y$ by using Bayes' theorem to combine observed and prior information

$$p(x|y) = \frac{p(y|x)p(x)}{\int_{\mathbb{R}^d} p(y|\tilde{x})p(\tilde{x})d\tilde{x}}$$

Latent Diffusion Models




Latent Diffusion scheme (Source NeurIPS 2023 Tutorial)



$$d\mathbf{z}_t = -\frac{\beta_t}{2}\mathbf{z}_t dt + \sqrt{\beta_t}d\mathbf{w},$$
$$d\mathbf{z}_t = \left[-\frac{\beta_t}{2}\mathbf{z}_t - \beta_t \nabla_{\mathbf{z}_t} \log p_t(\mathbf{z}_t) \right] dt + \sqrt{\beta_t}d\mathbf{w},$$

$$\mathcal{E} : \mathbb{R}^n \mapsto \mathbb{R}^d, \quad \mathcal{D} : \mathbb{R}^d \mapsto \mathbb{R}^n, \quad \mathbf{x} \approx \mathcal{D}(\mathcal{E}(\mathbf{x})),$$

Probability Flow ODE & Consistency Models



Consistency Models:

A distilled diffusion model obtained by training a deep neural network to transport \mathbf{x}_t to \mathbf{x}_0 by mapping any point on the ODE's trajectory back to the origin. CMs are one-step samplers.

Posterior Sampling Overdamped Langevin diffusion

$$d\mathbf{x}_s = \nabla \log p(\mathbf{y}|\mathbf{x}_s)ds + \nabla \log p(\mathbf{x}_s|c)ds + \sqrt{2}d\mathbf{w}_s$$

Key observations:

- Converges exponentially fast to the posterior $p(\mathbf{x}|\mathbf{y}, \mathbf{c})$ as the time s increases.
- Modular structure with explicit likelihood (data fidelity) and regularisation terms.
- No need to embed likelihood within reverse SDE/ODE through approximations.
- How do we replace $\nabla \log p(\mathbf{x}_s|c)$ by a generative model, e.g., stable diffusion ?

Proposed discrete-time approximation

$$\begin{aligned}\mathbf{u} &= \mathbf{x}_k + \int_0^\delta \nabla \log p(\tilde{\mathbf{x}}_s | c) ds + \sqrt{2} d\mathbf{w}_s, \quad \tilde{\mathbf{x}}_0 = \mathbf{x}_k, \\ \mathbf{x}_{k+1} &= \mathbf{u} + \delta \nabla \log p(\mathbf{y} | \mathbf{x}_{k+1}),\end{aligned}$$


Main observations:

- The first line corresponds to a Langevin SDE targeting the prior $\mathbf{p}(\mathbf{x}|c)$.
 - It admits $\mathbf{p}(\mathbf{x}|c)$ as unique invariant distribution.
 - It contracts exponentially fast towards $\mathbf{p}(\mathbf{x}|c)$ as δ increases.
- The second line (implicit Euler) is equivalent to a so-called proximal step that can be solved exactly for many imaging problems.
- Key idea: replace the first line by a different Markov kernel that has similar properties.


Auto-Encoding Stable Diffusion

$$\mathfrak{E}_t : \quad \mathbf{z}_t | \mathbf{x} \sim \mathcal{N}(\sqrt{\alpha_t} \mathcal{E}(\mathbf{x}), (1 - \alpha_t) \text{Id}_d)$$

$$\mathfrak{D}_{t,c} : \quad \mathbf{x}' = \mathcal{D}(G_\theta(\mathbf{z}'_t, t, c))$$



Proposed Plug-and-Play Langevin scheme



```
for  $k = 1, \dots, N$  do
     $\epsilon \sim \mathcal{N}(0, \text{Id})$ 
     $\mathbf{z}_{t_k}^{(k)} \leftarrow \sqrt{\alpha_{t_k}} E(\mathbf{x}^{(k-1)}) + \sqrt{1 - \alpha_{t_k}} \epsilon$     ▷ Encode
     $\mathbf{u}^{(k)} \leftarrow \mathcal{D}(G_\theta(\mathbf{z}_{t_k}^{(k)}, t_k, c))$                                 ▷ Decode
     $\mathbf{x}^{(k)} \leftarrow \text{prox}_{\delta_k g_y}(\mathbf{u}^{(k)})$     ▷  $g_y : \mathbf{x} \mapsto -\log p(\mathbf{y}|\mathbf{x})$ 
end for
```

LATINO (LAtent consisTency INverse sOlver)

Prompt Optimisation Stochastic Approximation Projected Gradient

$$\hat{c}(\mathbf{y}) = \arg \max_{c \in \mathbb{R}^k} p(\mathbf{y} \mid c)$$

$$c_{m+1} = \Pi_C[c_m + \gamma_m \nabla_c \log p(\mathbf{y} \mid c_m)]$$

$$\begin{aligned}\nabla_c \log p(\mathbf{y} \mid c) &= \mathbb{E}_{\mathbf{x} \mid \mathbf{y}, c} [\nabla_c \log p(\mathbf{y}, \mathbf{x} \mid c)], \\ &= \mathbb{E}_{\mathbf{x} \mid \mathbf{y}, c} [\nabla_c \log p(\mathbf{x} \mid c)],\end{aligned}$$

$$\nabla_c \log p(\mathbf{y} \mid c_m) \approx \nabla_c \log p(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid c_m)$$


$$c_{m+1} = \Pi_C [c_m + \gamma_m \nabla_c \log p(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid c_m)]$$

Prompt Optimisation Stochastic Approximation Projected Gradient

```
for  $m = 1, \dots, M$  do
    for  $k = 1, \dots, N_m$  do                                ▷ LATINO
         $\epsilon \sim \mathcal{N}(0, \text{Id})$ 
         $\mathbf{z}_{t_k}^{(k)} \leftarrow \sqrt{\alpha_{t_k}} \mathcal{E}(\mathbf{x}^{(k-1)}) + \sqrt{1 - \alpha_{t_k}} \epsilon$ 
         $\mathbf{u}^{(k)} \leftarrow \mathcal{D}(G_\theta(\mathbf{z}_{t_k}^{(k)}, t_k, c_m))$ 
         $\mathbf{x}^{(k)} \leftarrow \text{prox}_{\delta_k g_y}(\mathbf{u}^{(k)})$ 
    end for
     $h(c_m) \leftarrow \nabla_c \log p(\mathbf{z}_{t_1}^{(1)}, \dots, \mathbf{z}_{t_{N_m}}^{(N_m)} | c_m)$ 
     $c_{m+1} = \Pi_C [c_m + \gamma_m h(c_m)]$                       ▷ SAPG
     $\mathbf{x}^{(0)} \leftarrow \mathbf{x}^{(N_m)}$                                 ▷ Carry state forward
```

LATINO-PRO (LAtent consisTency INverse sOlver with PRompt Optimisation)

Some Results



LATINO (8 NFEs) & LATINO-PRO (68 NFEs)

Visualisation of Prompt Optimisation



Editing: sample from $p(x|c)$ using constrained SAPG steps to enforce semantic constraints

A sample from $p(x|c)$ before and after 4 SAPG steps
to adjust prompt (semantics)



Warning! The
Internet is Biased



Open questions for adventurous NAs

- Asymptotic and non-asymptotic convergence analysis for large δ .
- What Markov kernels are “good” approximations of $\mathbf{u} = \mathbf{x}_k + \int_0^\delta \nabla \log p(\tilde{\mathbf{x}}_s | c) ds + \sqrt{2} d\mathbf{w}_s$, $\tilde{\mathbf{x}}_0 = \mathbf{x}_k$, ?
- Constraining models to remain log-concave leads to worse models, but they also lead to slower algorithms. Why?
- No other known Langevin sampler (excluding trivial cases) converges in 4-8 steps in dimension 1M. We observe this behaviour with other DM priors, and on pixel space too. What’s going on here?
- Good strategies for moving the forward model to the latent space (save encoder-decoder evals.)

Thank you!
<https://arxiv.org/abs/2503.12615>